

Final Project: Practical Machine Learning

Angie

14 10 2019

```
library(caret)
```

```
## Loading required package: lattice
```

```
## Loading required package: ggplot2
```

```
library(rpart)
library(rpart.plot)
library(RColorBrewer)
library(rattle)
```

```
## Rattle: A free graphical interface for data science with R.
## Versi3n 5.2.0 Copyright (c) 2006-2018 Togaware Pty Ltd.
## Escriba 'rattle()' para agitar, sacudir y rotar sus datos.
```

```
library(randomForest)
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
##
## Attaching package: 'randomForest'
```

```
## The following object is masked from 'package:rattle':
##
## importance
```

```
## The following object is masked from 'package:ggplot2':
##
## margin
```

```
library(e1071)
library(gbm)
```

```
## Loaded gbm 2.1.5
```

Goal

The goal of your project is to predict the manner in which they did the exercise. This is the “classe” variable in the training set. You may use any of the other variables to predict with. You should create a report describing how you built your model, how you used cross validation, what you think the expected out of sample error is, and why you made the choices you did. You will also use your prediction model to predict 20 different test cases.

Setting data

```
set.seed(12345)

trainUrl <- "http://d396qusza40orc.cloudfront.net/predmachlearn/pml-training.csv"
testUrl <- "http://d396qusza40orc.cloudfront.net/predmachlearn/pml-testing.csv"

training <- read.csv(url(trainUrl), na.strings=c("NA", "#DIV/0!", ""))
testing <- read.csv(url(testUrl), na.strings=c("NA", "#DIV/0!", ""))
```

Setting Data

```
inTrain <- createDataPartition(training$classe, p=0.6, list=FALSE)
myTraining <- training[inTrain, ]
myTesting <- training[-inTrain, ]
dim(myTraining); dim(myTesting)
```

```
## [1] 11776 160
```

```
## [1] 7846 160
```

Remove Zero Variance variables

```
nzv <- nearZeroVar(myTraining, saveMetrics=TRUE)
myTraining <- myTraining[,nzv$nzv==FALSE]

nzv <- nearZeroVar(myTesting, saveMetrics=TRUE)
myTesting <- myTesting[,nzv$nzv==FALSE]
myTraining <- myTraining[c(-1)]
```

Pre-processing Data

```

trainingV3 <- myTraining
for(i in 1:length(myTraining)) {
  if( sum( is.na( myTraining[, i] ) ) /nrow(myTraining) >= .7) {
    for(j in 1:length(trainingV3)) {
      if( length( grep(names(myTraining[i]), names(trainingV3)[j]) ) == 1) {
        trainingV3 <- trainingV3[ , -j]
      }
    }
  }
}

# Set back to the original variable name
myTraining <- trainingV3
rm(trainingV3)

```

Processing Data Sets

```

clean1 <- colnames(myTraining)
clean2 <- colnames(myTraining[, -58]) # remove the classe column
myTesting <- myTesting[clean1]        # allow only variables in myTesting that are also in myTraining
testing <- testing[clean2]            # allow only variables in testing that are also in myTraining

dim(myTesting)

```

```
## [1] 7846 58
```

```
dim(testing)
```

```
## [1] 20 57
```

Format Data sets

```

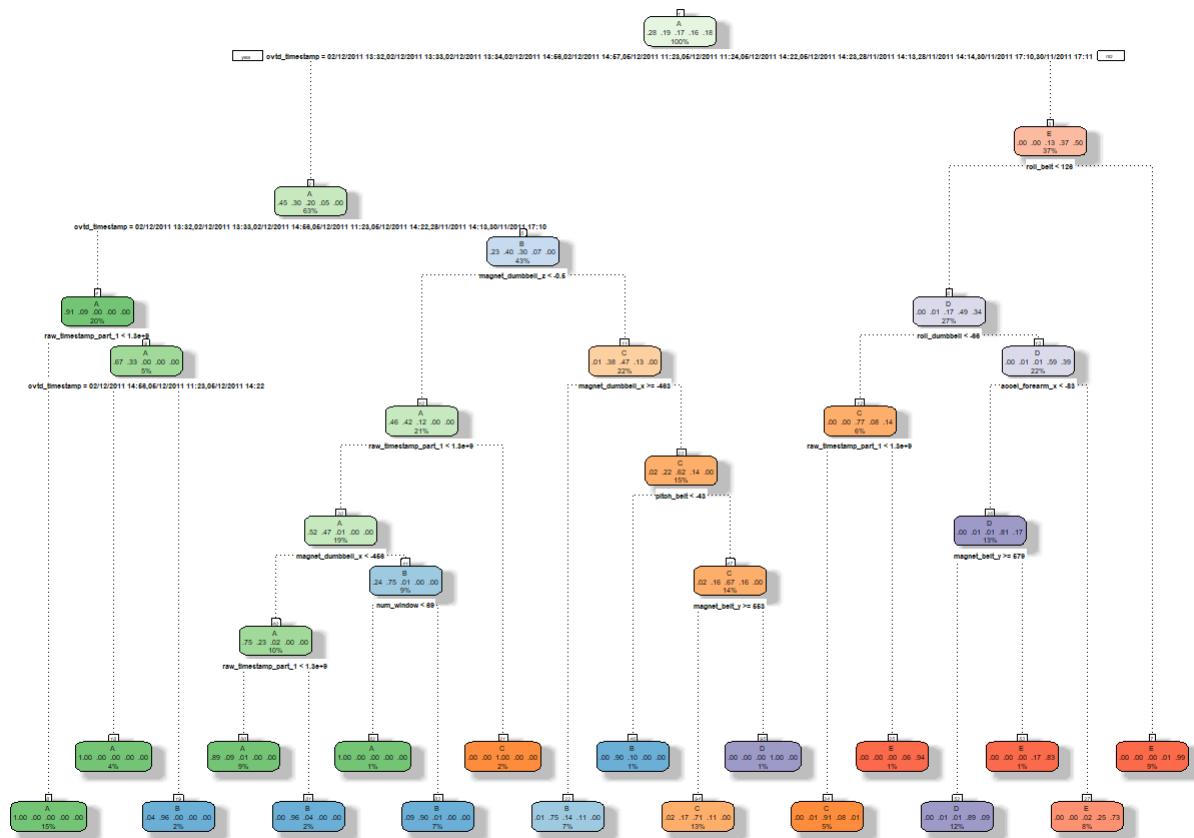
for (i in 1:length(testing) ) {
  for(j in 1:length(myTraining)) {
    if( length( grep(names(myTraining[i]), names(testing)[j]) ) == 1) {
      class(testing[j]) <- class(myTraining[i])
    }
  }
}

testing <- rbind(myTraining[2, -58] , testing)
testing <- testing[-1,]

```

Prediction 1: Decision Trees

```
set.seed(12345)
modFitA1 <- rpart(classe ~ ., data=myTraining, method="class")
fancyRpartPlot(modFitA1)
```



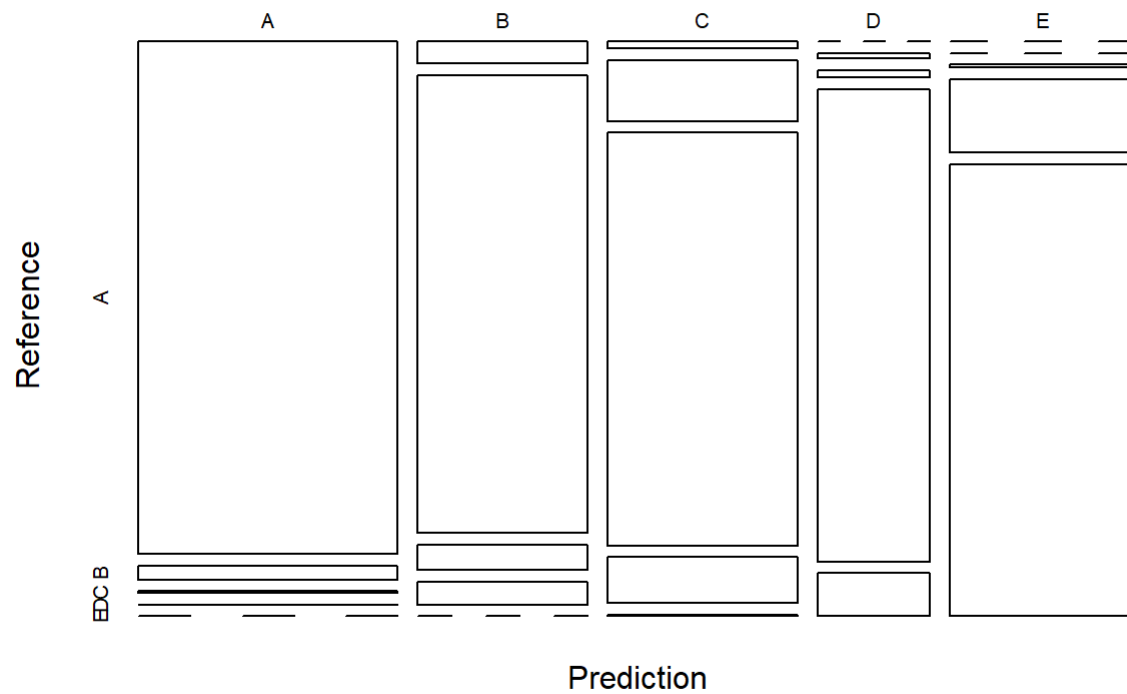
Rattle 2019-oct.-14 12:53:58 cmet10

```
predictionsA1 <- predict(modFitA1, myTesting, type = "class")
cmtree <- confusionMatrix(predictionsA1, myTesting$classe)
cmtree
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 2150   60    7    1    0
##           B   61 1260   69   64    0
##           C   21  188 1269  143    4
##           D    0   10   14  857   78
##           E    0    0    9  221 1360
##
## Overall Statistics
##
##           Accuracy : 0.8789
##           95% CI : (0.8715, 0.8861)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.8468
##
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9633  0.8300  0.9276  0.6664  0.9431
## Specificity      0.9879  0.9693  0.9450  0.9845  0.9641
## Pos Pred Value   0.9693  0.8666  0.7809  0.8936  0.8553
## Neg Pred Value   0.9854  0.9596  0.9841  0.9377  0.9869
## Prevalence       0.2845  0.1935  0.1744  0.1639  0.1838
## Detection Rate   0.2740  0.1606  0.1617  0.1092  0.1733
## Detection Prevalence 0.2827  0.1853  0.2071  0.1222  0.2027
## Balanced Accuracy 0.9756  0.8997  0.9363  0.8254  0.9536
```

```
plot(cmtree$table, col = cmtree$byClass, main = paste("Decision Tree Confusion Matrix: Accuracy
=", round(cmtree$overall['Accuracy'], 4)))
```

Decision Tree Confusion Matrix: Accuracy = 0.8789




Prediction 2: Random Forests

```
set.seed(12345)
modFitB1 <- randomForest(classe ~ ., data=myTraining)
predictionB1 <- predict(modFitB1, myTesting, type = "class")
cmrf <- confusionMatrix(predictionB1, myTesting$classe)
cmrf
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 2231    2    0    0    0
##           B    1 1516    1    0    0
##           C    0    0 1366    3    0
##           D    0    0    1 1281    1
##           E    0    0    0    2 1441
##
## Overall Statistics
##
##           Accuracy : 0.9986
##           95% CI : (0.9975, 0.9993)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9982
##
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9996  0.9987  0.9985  0.9961  0.9993
## Specificity      0.9996  0.9997  0.9995  0.9997  0.9997
## Pos Pred Value   0.9991  0.9987  0.9978  0.9984  0.9986
## Neg Pred Value   0.9998  0.9997  0.9997  0.9992  0.9998
## Prevalence       0.2845  0.1935  0.1744  0.1639  0.1838
## Detection Rate   0.2843  0.1932  0.1741  0.1633  0.1837
## Detection Prevalence 0.2846  0.1935  0.1745  0.1635  0.1839
## Balanced Accuracy 0.9996  0.9992  0.9990  0.9979  0.9995
```

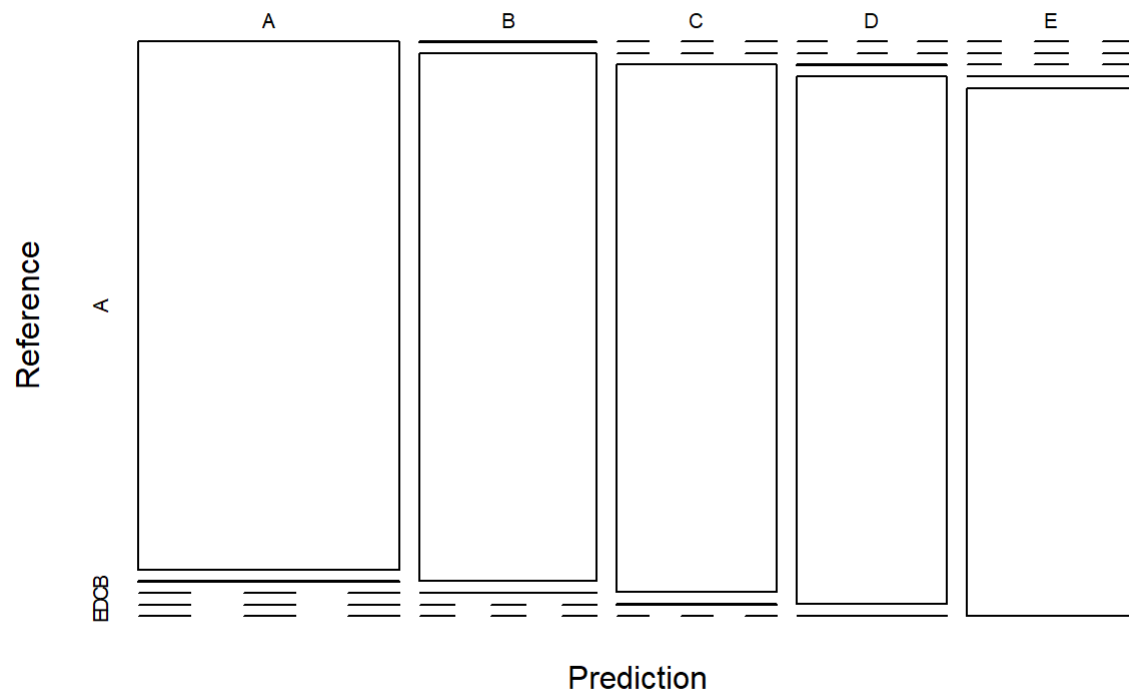
```
plot(modFitB1)
```

modFitB1



```
plot(cmrfr$table, col = cmtree$byClass, main = paste("Random Forest Confusion Matrix: Accuracy =",  
  , round(cmrfr$overall['Accuracy'], 4)))
```


Random Forest Confusion Matrix: Accuracy = 0.9986



Prediction 3: Generalized Boosted

```
set.seed(12345)
fitControl <- trainControl(method = "repeatedcv",
                           number = 5,
                           repeats = 1)


gbmFit1 <- train(classe ~ ., data=myTraining, method = "gbm",
                trControl = fitControl,
                verbose = FALSE)

gbmFinMod1 <- gbmFit1$finalModel

gbmPredTest <- predict(gbmFit1, newdata=myTesting)
gbmAccuracyTest <- confusionMatrix(gbmPredTest, myTesting$classe)
gbmAccuracyTest
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   A    B    C    D    E
##           A 2231    5    0    0    0
##           B    1 1512    0    0    0
##           C    0    1 1362    5    0
##           D    0    0    6 1272    0
##           E    0    0    0    9 1442
##
## Overall Statistics
##
##           Accuracy : 0.9966
##           95% CI : (0.995, 0.9977)
##           No Information Rate : 0.2845
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.9956
##
##           McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: A Class: B Class: C Class: D Class: E
## Sensitivity      0.9996  0.9960  0.9956  0.9891  1.0000
## Specificity      0.9991  0.9998  0.9991  0.9991  0.9986
## Pos Pred Value   0.9978  0.9993  0.9956  0.9953  0.9938
## Neg Pred Value    0.9998  0.9991  0.9991  0.9979  1.0000
## Prevalence       0.2845  0.1935  0.1744  0.1639  0.1838
## Detection Rate   0.2843  0.1927  0.1736  0.1621  0.1838
## Detection Prevalence 0.2850  0.1928  0.1744  0.1629  0.1849
## Balanced Accuracy 0.9993  0.9979  0.9973  0.9941  0.9993
```

```
plot(gbmFit1, ylim = c(0.9,1))
```



Results

The accuracy obtained from Random Forests is 99.89%. The expected sample error is $100 - 99.89 = 0.11\%$.

```
predictionB2 <- predict(modFitB1, testing, type = "class")
predictionB2
```

```
##  1  2 31  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19 20
##  B  A  B  A  A  E  D  B  A  A  B  C  B  A  E  E  A  B  B  B
## Levels: A B C D E
```